




NAZARBAYEV  
UNIVERSITY

SCHOOL OF SCIENCE AND TECHNOLOGY



# **ROBT305 - Embedded Systems**

**Lecture 4 – POSIX Pthreads Library and  
Intertask Communication and Process  
Synchronization: Mutexes**

**1-3 September, 2015**

# Course Logistics

---

## **Reading Assignment:**

**Chapters 3, 4, 6** of the of the Operating Systems Concept textbook (relevant material only)

**Chapter 3** of the Real-Time Systems Design and Analysis textbook (relevant material only)

**Homework Assignment #1** is out in Moodle and due to end of 13 September (Sunday)

**Quiz #1** is on 8 September – Process, Threads, Linux

---


# Process vs. Threads

---

**A process** (synonymously called “task”) is

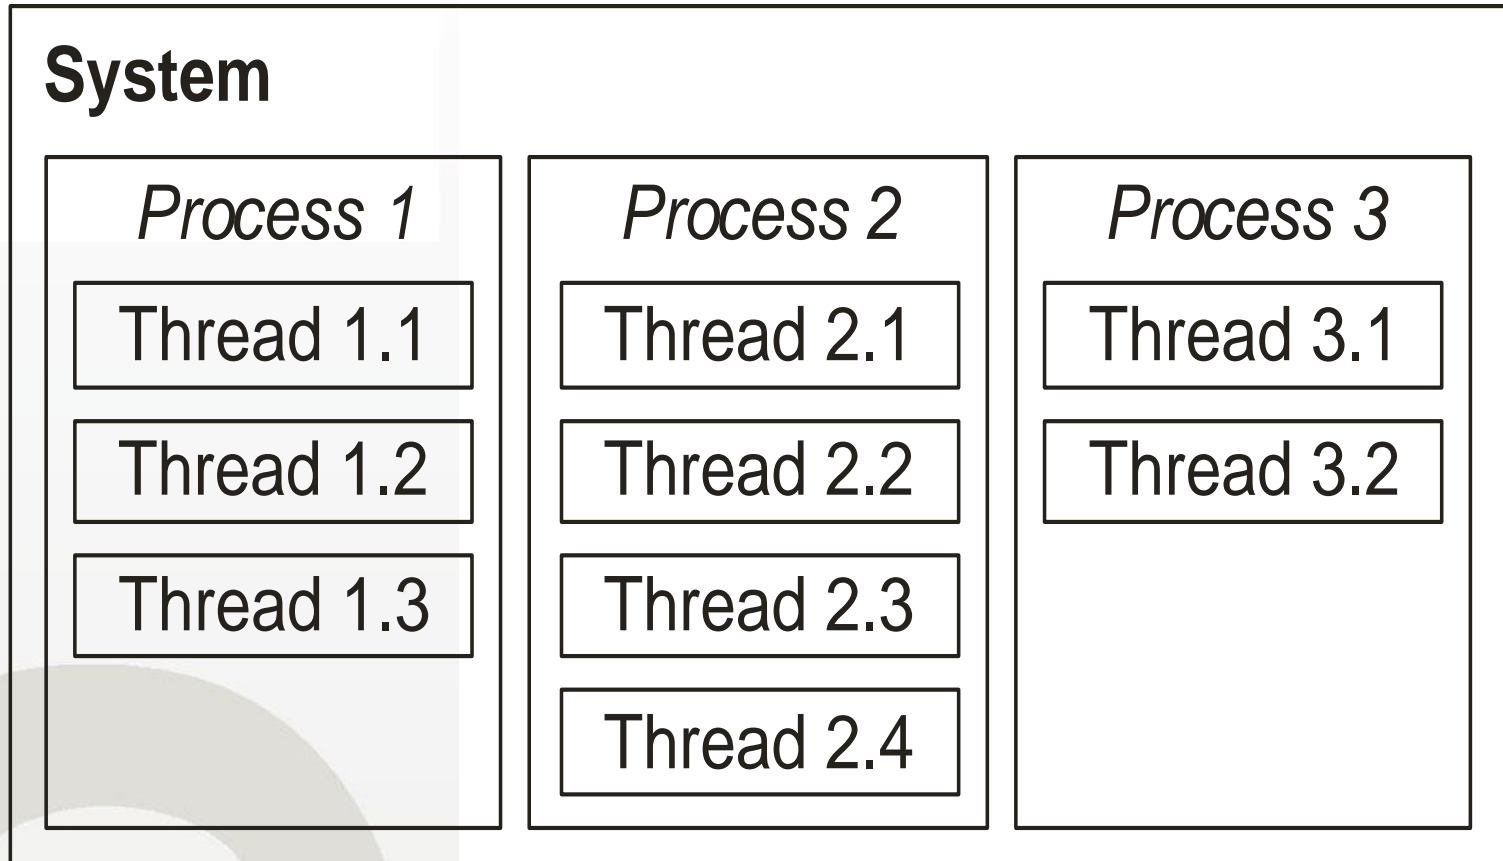
- An abstraction of a running program
- The logical unit of work schedulable by the OS

**A thread** is a lightweight process within a regular process. It has access to the same memory space, and the context switching from one thread to another will be shorter than for process to process



# Processes, and Multiple Threads

---



# POSIX

---

- ▶ **POSIX** (or "Portable Operating System Interface" is the collective name of a family of related standards specified by the IEEE to define the Application Programming Interface (API) for software compatible with variants of the Unix operating system.
- ▶ Originally, the name stood for IEEE Std 1003.1-1988. The family of POSIX standards is formally designated as **IEEE 1003** and the international standard name is ISO/IEC 9945. The standards emerged from a project that began in 1985. The term *POSIX* was suggested by Richard Stallman in response to an IEEE request for a memorable name!
- ▶ So, POSIX is not an OS, but an API! The API include Real-Time Services, Threads interface, Real-Time extensions, etc.

# Thread Libraries

---

- ▶ **Thread library** provides programmer with API for creating and managing threads
  - ▶ Two primary ways of implementing
    - ▶ Library entirely in user space
    - ▶ Kernel-level library supported by the OS
  - ▶ when compiling under gcc & GNU/Linux, remember the **-lpthread** option!
-

# POSIX Threads Programming

---

- ▶ POSIX Pthreads library
- ▶ May be provided either as user-level or kernel-level
- ▶ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- ▶ API specifies behavior of the thread library, implementation is up to development of the library
- ▶ Common in UNIX operating systems (Solaris, Linux, Mac OS X)

# Thread Body

---

- ▶ a thread is identified by a C function, called body:

```
void *my_thread(void *arg)
{
    ...
}
```

- ▶ a thread starts with the first instruction of its body
  - ▶ the thread ends when the body function ends
  - ▶ it's not the only way a thread can finish
-



# Thread Creation

---

- ▶ thread can be created using the primitive  

```
int pthread_create( pthread_t *ID,  
                  pthread_attr_t *attr,  
                  void *(*body)(void *),  
                  void * arg);
```
  - ▶ `pthread_t` is the type that contains the thread ID
  - ▶ `pthread_attr_t` is the type that contains the parameters of the thread
  - ▶ `arg` is the argument passed to the thread body when it starts
-

# Thread Termination

---

- ▶ a thread can terminate itself by calling  
`void pthread_exit(void *retval);`
  - ▶ when the thread body ends after the last “}”,  
`pthread_exit()` is called implicitly
  - ▶ exception: when `main()` terminates, `exit()` is called  
implicitly
-

# Thread IDs

---

- ▶ each thread has a unique ID
  - ▶ the thread ID of the current thread can be obtained using  
`pthread_t pthread_self(void);`
  - ▶ two thread IDs can be compared using  
`int pthread_equal( pthread_t thread1,  
pthread_t thread2 );`
-

# Joining a Thread

---

- ▶ a thread can wait the termination of another thread using  

```
int pthread_join( pthread_t th,  
                 void **thread_return);
```
  - ▶ it gets the return value of the thread or `PTHREAD_CANCELED` if the thread has been killed
  - ▶ by default, every task **must** be joined
  - ▶ the join frees all the internal resources (stack, registers, and so on)
-

# Pthreads Example

---

```
/*
 * DESCRIPTION: * A "hello world" Pthreads program. Demonstrates thread
 * creation and * termination. *
 */
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void *PrintHello(void *threadID)
{
    long tID;
    tID = (long) threadID;
    printf ("Hello World! It's me, thread #%ID!\n", tID);
    pthread_exit(NULL);
}
```

---

# Pthreads Example Cont.

---

```
int main (int argc, char *argv[])
{
    pthread_t  threads[NUM_THREADS];
    int rc;
    long t;
    for(t=0; t<NUM_THREADS; t++)
    {
        printf("In main: creating thread %ID\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc)
        {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    } /* Last thing that main() should do */
    pthread_exit(NULL);
}
```

---

# Pthreads Example

```
#include <pthread.h>
#include <stdio.h>

int sum; /* this data is shared by the thread(s) */
void *runner(void *param); /* threads call this function */

int main(int argc, char *argv[])
{
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if (argc != 2) {
        fprintf(stderr, "usage: a.out <integer value>\n");
        return -1;
    }
    if (atoi(argv[1]) < 0) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }
}
```

# Pthreads Example (Cont.)

```
/* get the default attributes */
pthread_attr_init(&attr);
/* create the thread */
pthread_create(&tid,&attr,runner,argv[1]);
/* wait for the thread to exit */
pthread_join(tid,NULL);

printf("sum = %d\n",sum);
}

/* The thread will begin control in this function */
void *runner(void *param)
{
    int i, upper = atoi(param);
    sum = 0;

    for (i = 1; i <= upper; i++)
        sum += i;

    pthread_exit(0);
}
```



# Pthreads Code for Joining 10 Threads

---

```
#define NUM_THREADS 10

/* an array of threads to be joined upon */
pthread_t workers[NUM_THREADS];

for (int i = 0; i < NUM_THREADS; i++)
    pthread_join(workers[i], NULL);
```

# Thread Cancellation

---

- ▶ Terminating a thread before it has finished
- ▶ Thread to be canceled is **target thread**
- ▶ Pthread code to create and cancel a thread:

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
. . .  
  
/* cancel the thread */  
pthread_cancel(tid);
```

---

# POSIX Threads Programming Practice

---

- ▶ Have a look at the beginning of this tutorial from Lawrence Livermore National Laboratory  
<https://computing.llnl.gov/tutorials/pthreads>
- ▶ A PDF version of the tutorial is on the Moodle web page

# Interprocess Communication

---

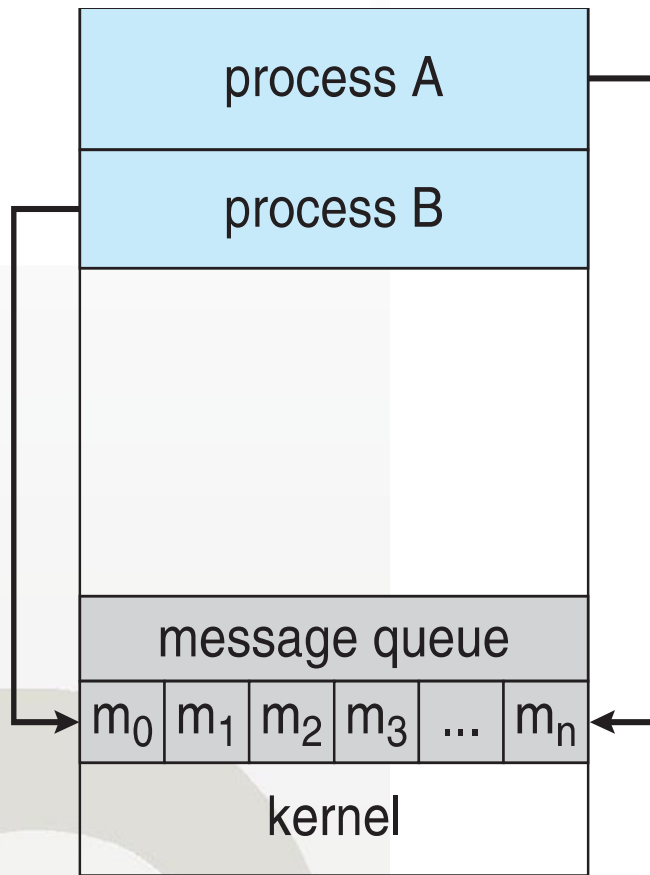
- ▶ Processes within a system may be **independent** or **cooperating**
  - ▶ Cooperating process can affect or be affected by other processes, including sharing data
  - ▶ Reasons for cooperating processes:
    - ▶ Information sharing
    - ▶ Computation speedup
    - ▶ Modularity
    - ▶ Convenience
  - ▶ Cooperating processes need **interprocess communication (IPC)**
-

# Intertask Communication

---


- ▶ Two models of intertask (interprocess) communication
    - ▶ Shared memory
    - ▶ Message passing
  - ▶ Task synchronization
    - ▶ Critical Regions
    - ▶ Mutexes
    - ▶ Semaphores
-

# Communications Models



(a)

Message Passing



(b)

Shared Memory

# Mailboxes (Messages)

---

- ▶ Messages are sent to and received from **mailboxes**, which provide an intertask communication mechanism
- ▶ **Mailbox** is a special memory location that one or more tasks can use to pass data, or more generally for synchronization.
- ▶ Tasks rely on the kernel to allow them to write to the location via a `post (send)` operation or to read from it via a `pend (receive)` operation.
- ▶ Mailboxes are available in most commercial RTOS.

# Buffering Data (Shared Memory)

---

- ▶ Mechanisms are needed to pass data between tasks in a multitasking system when production and consumption rates are unequal.
- ▶ Global variables are simple and fast, but have collision potential.
- ▶ Example, one task may produce data at a constant 1000 units per second, whereas another may consume these data at a rate less than 1000 units per second.
  - ▶ Assuming that the production interval is finite (and relatively short), the slower consumption rate can be accommodated if the producer fills a storage buffer with the data.
  - ▶ The buffer holds the excess data until the consumer task can catch up.
  - ▶ The buffer can be a queue or other data structure, including an unorganized mass of variables.
  - ▶ If consumer task consumes this information faster than it can be produced, or if the consumer cannot keep up with the producer, problems occur.
- ▶ Selection of the appropriate size buffer and synchronization mechanisms is critical in reducing or eliminating these problems.



# Time-Related Buffering

---

- ▶ Can use global variables for double buffering or ping-pong buffering.
- ▶ Used when time-related (correlated) data need to be transferred between cycles of different rates, or when a full set of data is needed by one process but can only be supplied slowly by another process.
- ▶ Variant of the classic bounded buffer problem in which a block of memory is used as a repository for data produced by “writers” and consumed by “readers.”
- ▶ Further generalization is the readers and writers problem in which there are multiple readers and multiple writers of a shared resource

# Time-Correlating Buffering Example

- ❑ IMU data measures  $x$ ,  $y$  and  $z$  accelerometer pulses in a 10 msec task
- ❑ Raw data processing is processed in the 40 msec task with lower priority (RM scheduling)
- ❑ Accelerometer data must be time-correlated, i.e., it is not allowed to process  $x(k)$  and  $y(k)$  with  $z(k+1)$  data samples – one reason for happening of such scenario can be the 10 msec task interrupt before  $z$  data processing.
- ❑ Solution – use buffering variables  $xb$ ,  $yb$  and  $zb$  in the 40 msec task with interrupt disabled.

```
int roff();          /* disable interrupts */
xb=x;               /* buffer x data */
yb=y;               /* buffer y data */
zb=z;               /* buffer z data */
int ron();           /* enable interrupts */
process(xb,yb,zb);  /* use buffered data */
...
```


# Ring Buffers

---

- ▶ A circular queue or ring buffer can be used to solve the problem of synchronizing multiple reader and writer tasks.
- ▶ Simultaneous input and output to the list are achieved by keeping head and tail indices.
- ▶ Data are loaded at the tail and read from the head.
- ▶

# Ring Buffers

- Tasks write to the buffer at the `tail` index and read data from the `head` index.
- Buffer is implemented as a circular array with two pointers `in(tail)` and `out(head)`
- Data access is synchronized with a counting semaphore set to the size of ring buffer `BUFFER_SIZE (n)`.
- Buffer is empty when:  
 $\text{tail}(\text{in}) == \text{head}(\text{out})$
- Buffer is full when:  
 $((\text{in} + 1) \% \text{BUFFER\_SIZE}) == \text{out}$



# Ring Buffering Example

---

- Suppose the ring buffer is a data structure of type `ring_buffer`

```
typedef struct ring_buffer
{
    int contents[n]; /* buffer area */
    int head=0;      /* head index */
    int tail=0;      /* tail index */
}
```

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

# Ring Buffering Example Cont

- Implementation of the `write(data, &s)` and `read(data, &s)` operations, writing and reading data to the ring buffer `s` respectively.

```
void write(int data, ring_buffer *s)
{
    if ((s->tail+1) % n == head)
        error();    /* buffer overflow */
    else
    {
        s->contents+tail=data;    /* write data */
        tail=(tail+1) % n;        /* update tail */
    }
}
```

```
    item next_produced;
```

```
    while (true) {
```

```
        /* produce an item in next_produced */
```

```
        while (((in + 1) % BUFFER_SIZE) == out)
```

```
            ; /* do nothing */
```

```
        buffer[in] = next_produced;
```

```
        in = (in + 1) % BUFFER_SIZE;
```

```
    }
```

# Ring Buffering Example Cont

```
void read(int data, ring_buffer *s)
{
    if (s->head==s->tail)
        data=NULL; /* buffer underflow */
    else
    {
        data=s->contents+head; /* read data */
        s->head=(s->head+1) % n; /* update head */
    }
}
```

Allows at most  
BUFFER\_SIZE-1 items in  
the buffer at the same time

```
item next_consumed;
while (true) {
    while (in == out)
        ; /* do nothing */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume the item in next consumed */
}
```

# Ring Buffering Example Cont

- Suppose that we need to fill **all** the buffer `BUFFER_SIZE`.
- Integer **counter** keeps track of the number of full buffer.
- Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.

```
while (true) {  
    /* produce an item in next produced */  
    while (counter == BUFFER_SIZE) ;  
        /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
    counter++;  
}
```

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    /* consume the item in next consumed */  
}
```



# Race Condition

- ▶ `counter++` could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

- ▶ `counter--` could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

To guard against the **race condition**, only one task at a time can be manipulating the variable **counter**. **Tasks should be synchronized**

- ▶ Consider this execution interleaving with “count = 5” initially:

S0: producer execute `register1 = counter`

{register1 = 5}

S1: producer execute `register1 = register1 + 1`

{register1 = 6}

S2: consumer execute `register2 = counter`

{register2 = 5}

S3: consumer execute `register2 = register2 - 1`

{register2 = 4}

S4: producer execute `counter = register1`

{counter = 6}

S5: consumer execute `counter = register2`

{counter = 4}

# Critical Regions

---

- ▶ Multitasking systems are concerned with resource sharing.
  - ▶ When resources can only be used by one task at a time, and use of the resource cannot be interrupted they are said to be “**serially reusable**”
  - ▶ Serial reusable devices include certain peripherals, shared memory, and the CPU.
  - ▶ The CPU protects itself against simultaneous use.
  - ▶ Code that interacts with the other serially reusable resources is called a “**critical region**” (“**critical section**”)
  - ▶ If two tasks enter the same critical region (section) simultaneously, a catastrophic error can occur.
  - ▶ Simultaneous use of a serial reusable resource results in a “**collision**.”
  - ▶ The concern is to provide a mechanism for preventing collisions.
-

# Critical Regions

---

- ▶ Consider two C programs, **Task\_A** and **Task\_B**, in a round-robin system.
- ▶ **Task\_B** outputs the message “**I am task\_B**” and **Task\_A** outputs the message “**I am Task\_A**.”
- ▶ In the midst of printing, Task\_B is interrupted by Task\_A, which begins printing. The result is the incorrect output:  
**I am I am Task\_A Task\_B**
- ▶ More serious complications could arise if both tasks were controlling devices in an embedded system.

# Critical Regions

---

- A task may have a **critical region (section)** segment of code changing common variables, updating table, writing file, etc.
- When one task in critical section, no other may be in its critical section
- Each task must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

---

# Synchronization Hardware

---

- ▶ Many systems provide hardware support for critical section code
  - ▶ All solutions below based on idea of **locking** - Protecting critical regions via locks
  - ▶ Single-processor environment – interrupts could be disabled
    - ▶ *Currently running critical region code execute without preemption*
    - ▶ *No unexpected modifications could be made to the shared variable*
  - ▶ Modern machines provide special atomic hardware instructions (system calls) to solve the critical region problem
    - ▶ **Atomic** = non-interruptible
-

# Mutex Locks

---

- ▶ **Mutex lock** - simplest software tool to solve critical region problems
    - ▶ The term *mutex* is short for *mutual exclusion*
  - ▶ A process must acquire the lock () before entering a critical region - **acquire()** function;
  - ▶ It releases the lock when it exits the critical region - **release()** function
    - ▶ A boolean variable **available** indicate if lock is available or not
  - ▶ Calls to **acquire()** and **release()** must be atomic
    - ▶ Usually implemented via hardware atomic instructions
  - ▶ But this solution requires **busy waiting**
    - ▶ This lock therefore called a **spinlock**
-

# acquire() and release()

---

```
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```

The definitions:

```
acquire()  
{  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}
```

```
release()  
{  
    available = true;  
}
```

---

# Mutex Lock Example

```
--- #define MAX_PROCESSES 255  
int number_of_processes = 0;  
  
/* the implementation of fork() calls this function */  
int allocate_process() {  
    int new_pid;  
  
    if (number_of_processes == MAX_PROCESSES)  
        return -1;  
    else {  
        /* allocate necessary process resources */  
        ++number_of_processes;  
  
        return new_pid;  
    }  
}  
  
/* the implementation of exit() calls this function */  
void release_process() {  
    /* release process resources */  
    --number_of_processes;  
}
```

- a. Identify the race condition(s).
- b. Assume you have a mutex lock named `mutex` with the operations `acquire()` and `release()`. Indicate where the locking needs to be placed to prevent the race condition(s).



# Mutex Lock Example

```
-----  
#define MAX_PROCESSES 255  
int number_of_processes = 0;  
  
/* the implementation of fork() calls this function */  
int allocate_process() {  
    int new_pid;  
        acquire() ;  
    if (number_of_processes == MAX_PROCESSES)  
        return -1;  
    else {  
        /* allocate necessary process resources */  
        ++number_of_processes;  
        release() ;  
        return new_pid;  
    }  
}  
  
/* the implementation of exit() calls this function */  
void release_process() {    acquire() ;  
    /* release process resources */  
    --number_of_processes;    release() ;  
}
```

- Identify the race condition(s).
- Assume you have a mutex lock named `mutex` with the operations `acquire()` and `release()`. Indicate where the locking needs to be placed to prevent the race condition(s).

# Pthreads Synchronization

---

- ▶ Mutex locks – main synchronization techniques used with Pthreads

```
#include <pthread.h>

pthread_mutex_t mutex;

/* create the mutex lock */
pthread_mutex_init(&mutex, NULL);

/* acquire the mutex lock */
pthread_mutex_lock(&mutex);

/* critical section */

/* release the mutex lock */
pthread_mutex_unlock(&mutex);
```

*If the mutex lock is unavailable when **thread\_mutex\_lock()** is invoked, the calling thread is blocked until the owner invokes **pthread\_mutex\_unlock()***

# Pthreads Synchronization Example

```
#include<stdio.h>
#include<string.h>
#include<pthread.h>
#include<stdlib.h>
#include<unistd.h>

pthread_t tid[2];
int counter;

void* doSomething(void *arg)
{
    unsigned long i = 0;
    counter += 1;
    printf("\n Job %d started\n", counter);

    for(i=0; i<(0xFFFFFFFF);i++);
    printf("\n Job %d finished\n", counter);

    return NULL;
}
```

```
int main(void)
{
    int i = 0;
    int err;

    while(i < 2)
    {
        err = pthread_create(&(tid[i]), NULL, &doSomething, NULL);
        if (err != 0)
            printf("\ncan't create thread :[%s]", strerror(err));
        i++;
    }

    pthread_join(tid[0], NULL);
    pthread_join(tid[1], NULL);

    return 0;
}
```

# Pthreads Synchronization Example

---

**Result of the program execution:**

```
$ ./tgsthreads  
Job 1 started  
Job 2 started  
Job 2 finished  
Job 2 finished
```

**Insert Mutexes to prevent race condition**

---

# Pthreads Synchronization Example

## Insert Mutexes to prevent race condition

```
#include<stdio.h>
#include<string.h>
#include<pthread.h>
#include<stdlib.h>
#include<unistd.h>

pthread_t tid[2];
int counter;
pthread_mutex_t lock;

void* doSomething(void *arg)
{
    pthread_mutex_lock(&lock);

    unsigned long i = 0;
    counter += 1;
    printf("\n Job %d started\n", counter);

    for(i=0; i<(0xFFFFFFFF);i++);

    printf("\n Job %d finished\n", counter);

    pthread_mutex_unlock(&lock);

    return NULL;
}
```

```
int main(void)
{
    int i = 0;
    int err;

    if (pthread_mutex_init(&lock, NULL) != 0)
    {
        printf("\n mutex init failed\n");
        return 1;
    }

    while(i < 2)
    {
        err = pthread_create(&(tid[i]), NULL, &doSomething, NULL);
        if (err != 0)
            printf("\ncan't create thread :[%s]", strerror(err));
        i++;
    }

    pthread_join(tid[0], NULL);
    pthread_join(tid[1], NULL);
    pthread_mutex_destroy(&lock);

    return 0;
}
```

# Any Questions?

---

