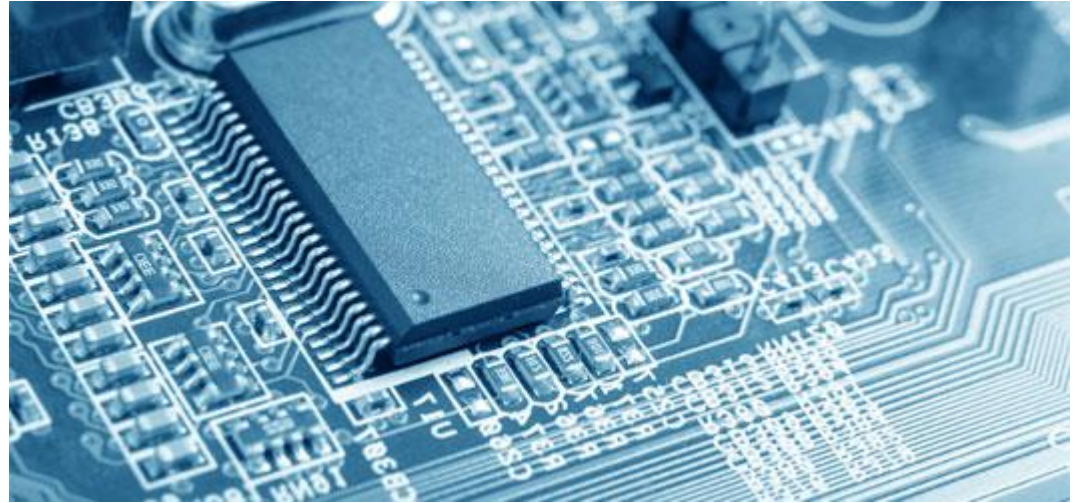




NAZARBAYEV
UNIVERSITY
SCHOOL OF SCIENCE AND TECHNOLOGY



ROBT305 - Embedded Systems

Lectures 10-11 – EDF Scheduling

**29 September -
6 October, 2015**

Course Logistics

Reading Assignment:

Chapter 5 of the Operating Systems Concept textbook

Chapter 3 of the Real-Time Systems Design and Analysis textbook

Homework Assignment #2 is out in Moodle and due to **end of 4 October (Sunday) – extended deadline**

Quiz #3 is on **1 October – Semaphores, Scheduling (except EDF)**

Midterm exam is in class time on Thursday 8 October, 2015

Dynamic Priority Scheduling

- ▶ In contrast to fixed-priority algorithms, in dynamic priority schemes the priority of the task with respect to that of the other tasks changes as tasks are released and completed.
- ▶ One of the most well-known dynamic algorithm, **earliest-deadline-first** (EDF), deals with deadlines rather than execution times.
- ▶ The ready task with the earliest deadline has the highest priority at any point of time.

Earliest Deadline First (EDF)

Theorem [EDF Bound]

A set of n periodic tasks, each of whose relative deadline equals its period, can be feasibly scheduled by EDF if and only if


$$\sum_{i=1}^n (e_i / p_i) \leq 1$$

EDF Example 2

Task set

| τ_i | e_i | p_i |
|----------|-------|-------|
| τ_1 | 2 | 5 |
| τ_2 | 4 | 7 |

EDF Example 2



EDF schedule for task set shown. Although τ_1 and τ_2 release simultaneously, τ_1 executes first because its deadline is earliest. At $t = 2$, τ_2 can execute. Even though τ_1 releases again at $t = 5$, its deadline is not earlier than τ_2 's. This sequence continues until time $t = 15$ when τ_2 is preempted as its deadline is later ($t = 21$) than τ_1 's ($t = 20$). τ_2 resumes when τ_1 completes.

EDF Example


| Task | Execution time, e_i | Period, p_i |
|-------|-----------------------|---------------|
| P_1 | 25 | 50 |
| P_2 | 35 | 80 |

- ▶ Failed to meet RM scheduling requirements
 - ▶ Priorities are assigned according to deadlines:
the earlier the deadline, the higher the priority;
the later the deadline, the lower the priority
-

EDF Example

| Task | Execution time, e_i | Period, p_i |
|-------|-----------------------|---------------|
| P_1 | 25 | 50 |
| P_2 | 35 | 80 |

- ▶ Failed to meet RM scheduling requirements
- ▶ Priorities are assigned according to deadlines:
the earlier the deadline, the higher the priority;
the later the deadline, the lower the priority



EDF Scheduling


Given are five tasks with arrival times, execution times and deadlines according to the following table.

| Task | Release time, r_i | Exec. Time, e_i | Deadline, d_i |
|------|---------------------|-------------------|-----------------|
| J1 | 0 | 3 | 16 |
| J2 | 2 | 1 | 7 |
| J3 | 0 | 6 | 8 |
| J4 | 8 | 2 | 11 |
| J5 | 13 | 3 | 18 |

Determine the EDF schedule.

EDF Scheduling

| Task | Release time, r_i | Exec. Time, e_i | Deadline, d_i |
|------|---------------------|-------------------|-----------------|
| J1 | 0 | 3 | 16 |
| J2 | 2 | 1 | 7 |
| J3 | 0 | 6 | 8 |
| J4 | 8 | 2 | 11 |
| J5 | 13 | 3 | 18 |



Earliest Deadline First


- ▶ EDF is more flexible and achieves better utilization than RM.
- ▶ However, the timing behavior of a system scheduled according to a fixed-priority algorithm is more predictable than that of a system scheduled according to a dynamic-priority-algorithm.
- ▶ In case of overloads, RM is stable in the presence of missed deadlines; the same lower priority tasks miss deadlines every time. There is no effect on higher priority tasks.
- ▶ When tasks are scheduled using EDF, it is difficult to predict which tasks will miss their deadlines during overloads.
- ▶ A good overrun management scheme is thus needed for such dynamic priority algorithms employed in systems where overload conditions cannot be avoided.

Problem of Mixed Task Sets

- ▶ In many applications, there are as well aperiodic as periodic tasks.
 - ▶ Periodic tasks: time-driven, execute critical control activities with hard timing constraints aimed at guaranteeing regular activation rates.
 - ▶ Aperiodic tasks: event-driven, may have hard, soft, nonreal-time requirements depending on the specific application.
 - ▶ Sporadic tasks: Offline guarantee of event-driven aperiodic tasks with critical timing constraints can be done only by making proper assumptions on the environment; that is by assuming a maximum arrival rate for each critical event.
-


Background Scheduling

- ▶ **Simple solution** for RM and EDF scheduling of periodic tasks:
 - Processing of aperiodic tasks in the background, i.e. if there are no periodic request.
 - Periodic tasks are not affected.
 - Response of aperiodic tasks may be prohibitively long and there is no possibility to assign a higher priority to them.



Background Scheduling

- **Example** (rate monotonic periodic schedule):



POSIX Upgrade 1b: Real-Time Extensions

- ▶ Priority Scheduling
- ▶ Real-Time Signals
- ▶ Clocks and Timers
- ▶ Semaphores
- ▶ Message Passing
- ▶ Shared Memory
- ▶ Asynch and Synch I/O
- ▶ Memory Locking
- ▶ Mostly compliant OS: Linux
- ▶ Some compliant: VxWorks RTOS

POSIX Real-Time Scheduling

- ❑ The POSIX.1b standard
 - ❑ API provides functions for managing real-time threads
 - ❑ Defines scheduling classes for real-time threads:
 1. **SCHED_FIFO** - threads are scheduled using a FCFS strategy with a FIFO queue. There is no time-slicing for threads of equal priority
 2. **SCHED_RR** - similar to SCHED_FIFO except time-slicing occurs for threads of equal priority
 3. **SCHED_OTHER** – undefined and system specific
 - ❑ Defines two functions for getting and setting scheduling policy:
 1. `pthread_attr_getsched_policy(pthread_attr_t *attr, int *policy)`
 2. `pthread_attr_setsched_policy(pthread_attr_t *attr, int policy)`
-

POSIX Real-Time Scheduling API

```
#include <pthread.h>
#include <stdio.h>
#define NUM_THREADS 5
int main(int argc, char *argv[])
{
    int i, policy;
    pthread_t tid[NUM_THREADS];
    pthread_attr_t attr;
    /* get the default attributes */
    pthread_attr_init(&attr);
    /* get the current scheduling policy */
    if (pthread_attr_getschedpolicy(&attr, &policy) != 0)
        fprintf(stderr, "Unable to get policy.\n");
    else {
        if (policy == SCHED_OTHER) printf("SCHED_OTHER\n");
        else if (policy == SCHED_RR) printf("SCHED_RR\n");
        else if (policy == SCHED_FIFO)
            printf("SCHED_FIFO\n");
    }
}
```

POSIX Real-Time Scheduling API (Cont.)

```
/* set the scheduling policy - FIFO, RR, or OTHER */
if (pthread_attr_setschedpolicy(&attr, SCHED_FIFO) != 0)
    fprintf(stderr, "Unable to set policy.\n");
/* create the threads */
for (i = 0; i < NUM_THREADS; i++)
    pthread_create(&tid[i], &attr, runner, NULL);
/* now join on each thread */
for (i = 0; i < NUM_THREADS; i++)
    pthread_join(tid[i], NULL);
}

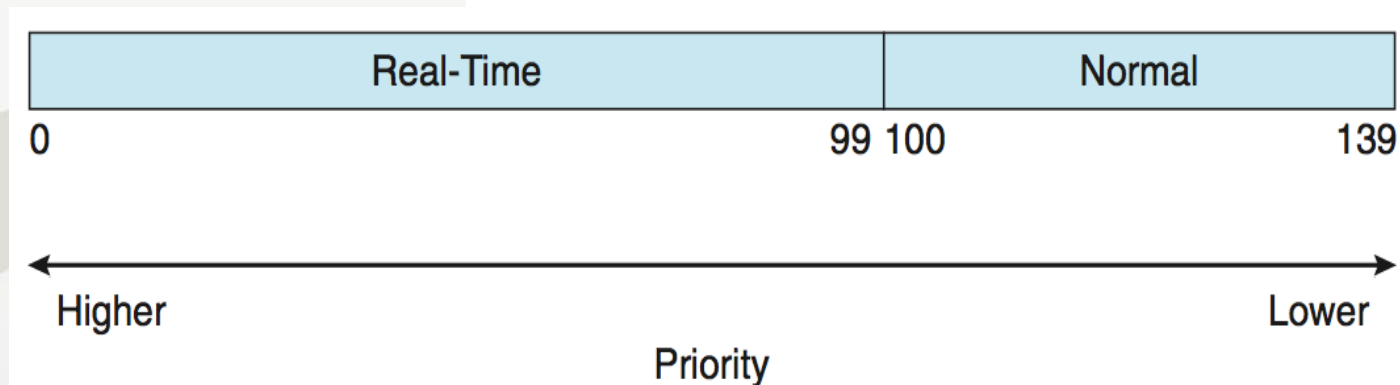
/* Each thread will begin control in this function */
void *runner(void *param)
{
    /* do some work ... */
    pthread_exit(0);
}
```

Linux Scheduling

- ▶ **Completely Fair Scheduler (CFS)**
 - ▶ **Scheduling classes**
 - ▶ Each has specific priority
 - ▶ Scheduler picks highest priority task in highest scheduling class
 - ▶ Rather than quantum based on fixed time allotments, based on proportion of CPU time
 - ▶ 2 scheduling classes included, others can be added
 1. default
 2. real-time
 - ▶ Quantum calculated based on **nice value** from -20 to +19
 - ▶ Lower value is higher priority
 - ▶ Calculates **target latency** – interval of time during which task should run at least once
 - ▶ Target latency can increase if say number of active tasks increases
 - ▶ CFS scheduler maintains per task **virtual run time**
 - ▶ Associated with decay factor based on priority of task – lower priority is higher decay rate
 - ▶ Normal default priority yields virtual run time = actual run time
 - ▶ To decide next task to run, scheduler picks task with lowest virtual run time
-

Linux Scheduling (Cont.)

- ▶ Real-time scheduling according to POSIX.1b
 - ▶ Real-time tasks have static priorities
- ▶ Real-time plus normal map into global priority scheme
- ▶ Nice value of -20 maps to global priority 100
- ▶ Nice value of +19 maps to priority 139



Any Questions?

